

CPSC 351 Operating System Concepts

Chapter 3: Processes

Chapter3-A

Chapter 3: Processes

- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication
- Examples of IPC Systems
- Communication in Client-Server Systems

Objectives

- To introduce the notion of **a process** -- a program in execution, which forms the basis of all computation
- To describe the various features of processes, including **scheduling, creation and termination, and communication**
- To explore **interprocess communication** using **shared memory** and **message passing**
- To describe communication in **client-server systems**

Process Review

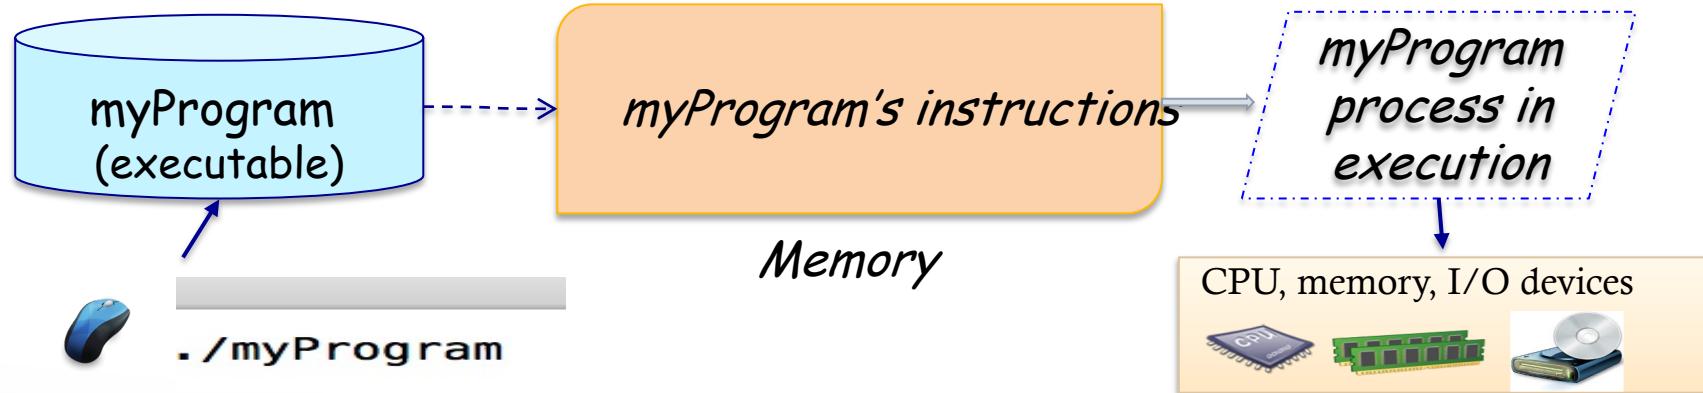
A *process* can be defined as:

A program in execution

An instance of a running program

The entity that can be assigned to, and executed on, a processor

A unit of activity characterized by a single sequential thread of execution, a current state, and an associated set of system resources



Process Concept

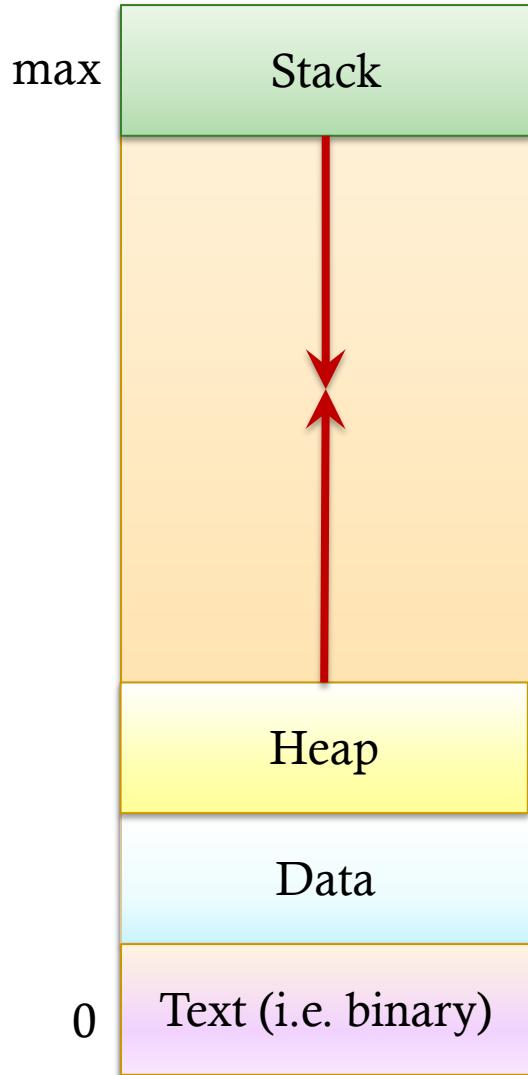
- An operating system executes a variety of programs:
 - Batch system - **jobs**
 - Time-shared systems - **user programs** or **tasks**
- Textbook uses the terms *job* and *process* almost interchangeably
- **Process** - a program in execution; process execution must progress in sequential fashion
- Multiple parts
 - ◆ The program code, also called **text section**
 - ◆ Current activity including **program counter**, processor registers
 - ◆ **Stack** containing temporary data
 - ▣ Function parameters, return addresses, local variables
 - ◆ **Data section** containing global variables
 - ◆ **Heap** containing memory dynamically allocated during run time

Process Concept (Cont.)

- Program is *passive* entity stored on disk (**executable file**), process is *active*
 - Program becomes process when executable file loaded into memory
- Execution of program started via **GUI mouse clicks, command line entry of its name, etc**
- One program can be several processes
 - Consider multiple users executing the same program

Process in Memory

- ❑ A process is more than the program code
- ❑ Includes
 - Current activity such as program counter (PC specifying the next instruction to execute) and contents of processor registers
 - **Stack:** stores function parameters, local variables, and return address
 - **Heap:** contains dynamic memory allocated during process runtime
 - **Data:** contains global variables
 - **Text:** stores the program instructions (i.e., the executable)



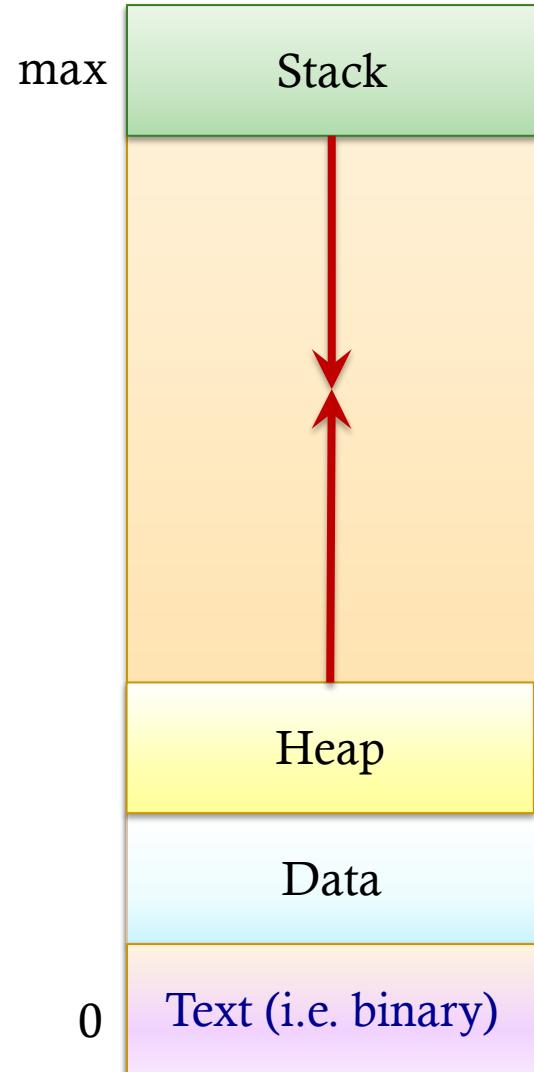
Process in Memory Example

```
int a = 1; Data
int f(int c) { int b = c+1; return b; }

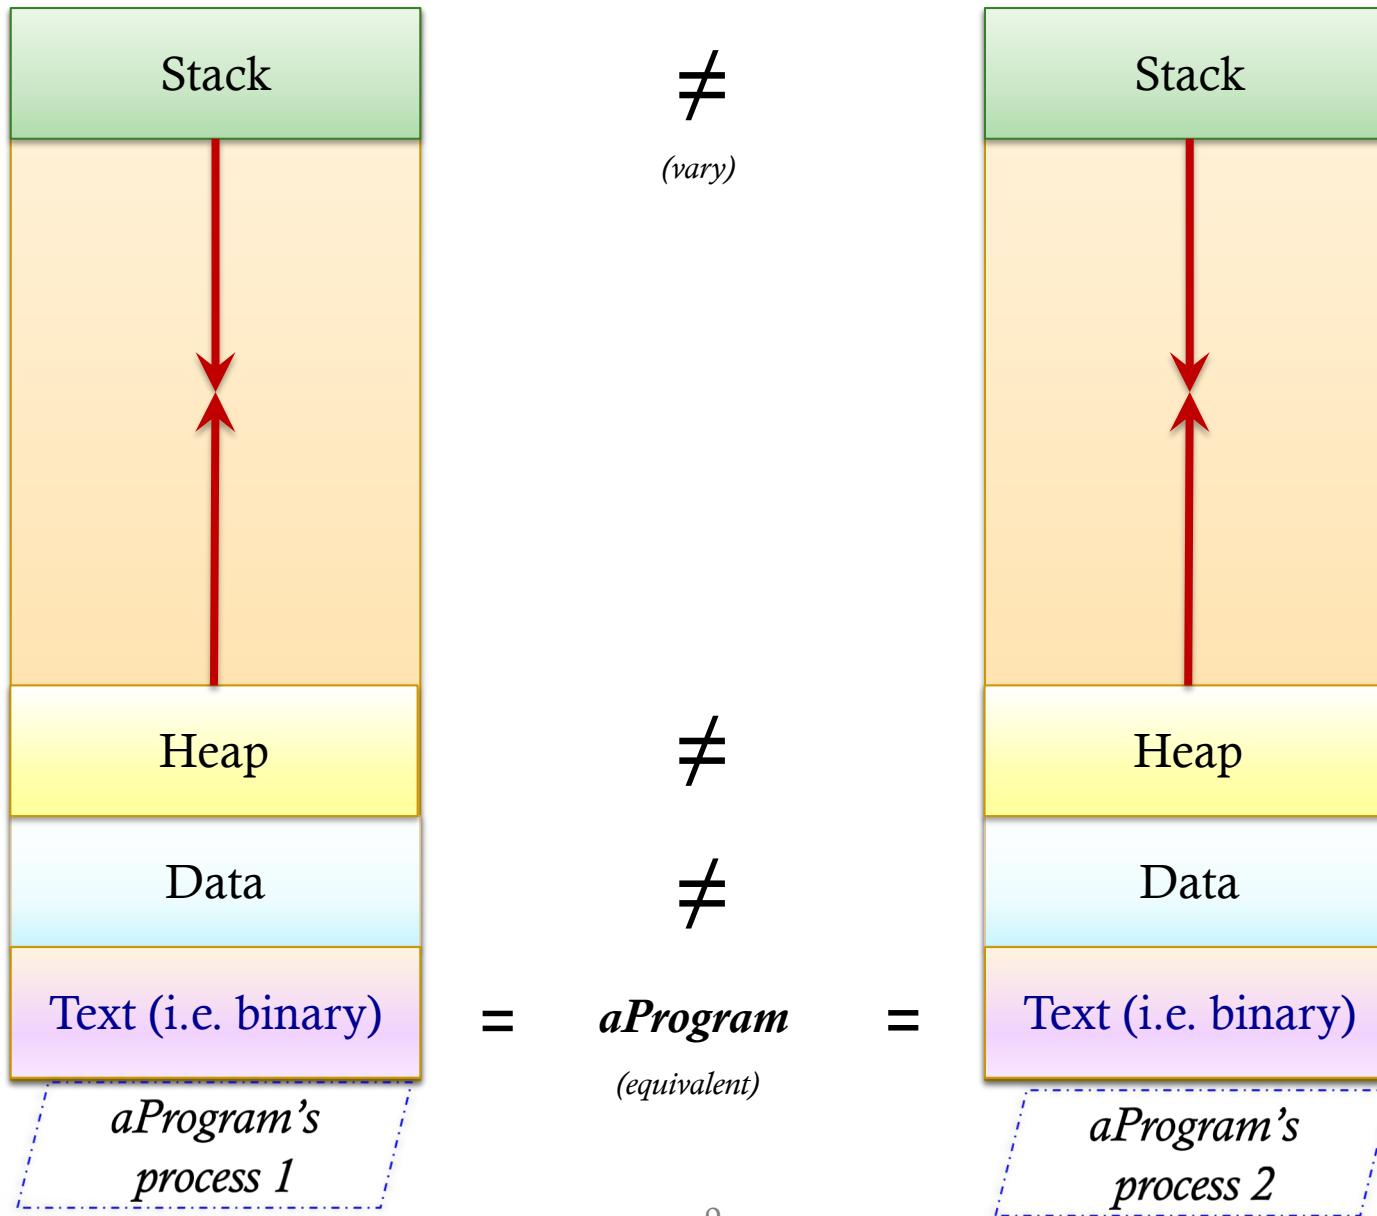
int main()
{
    int d = 2; Heap
    char* arr = new char[100];
    Stack for(int i = 0; i < 100; ++i)
        arr[i] = '\0';

    f(d);
    delete arr;
    return 0;
}
```

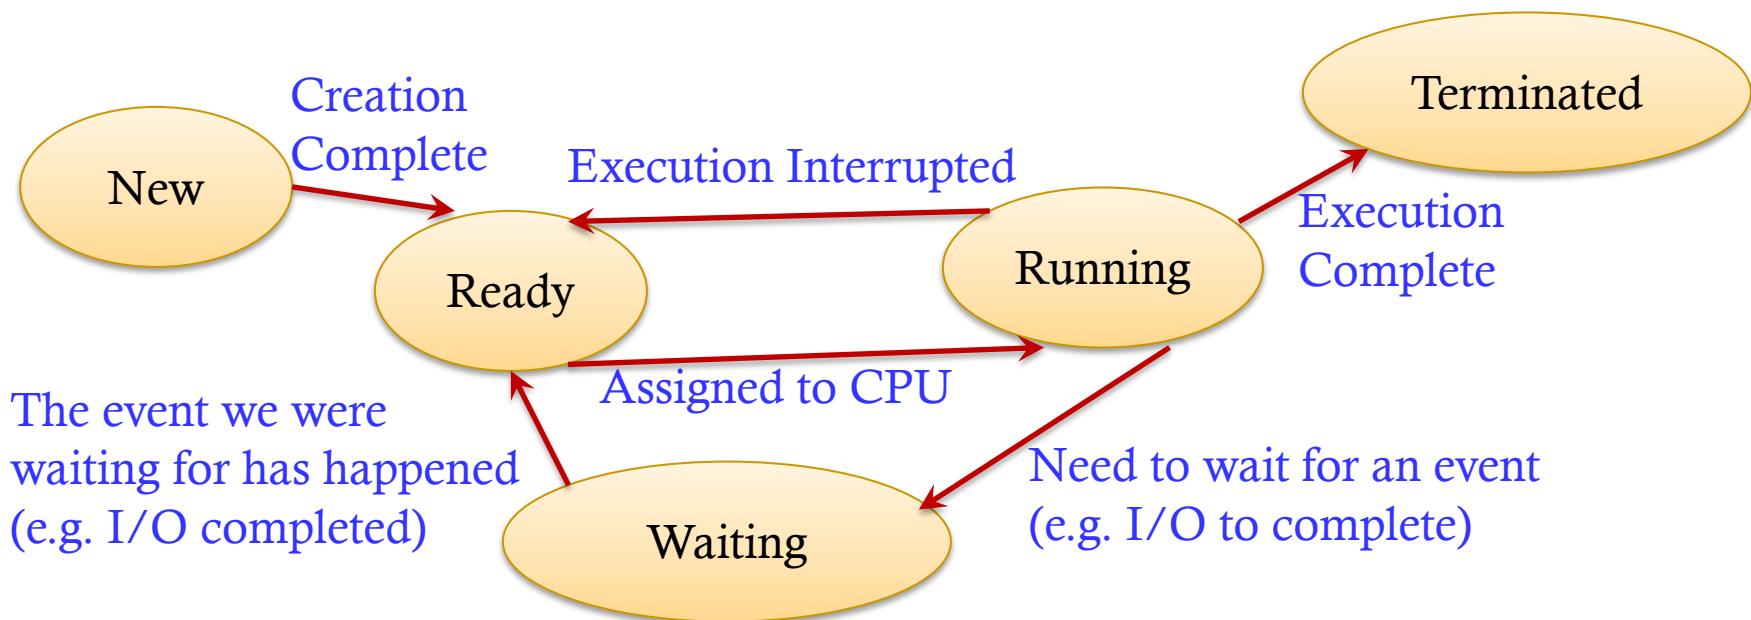
Text (i.e., executable instructions)



One Program Two Processes



Process State



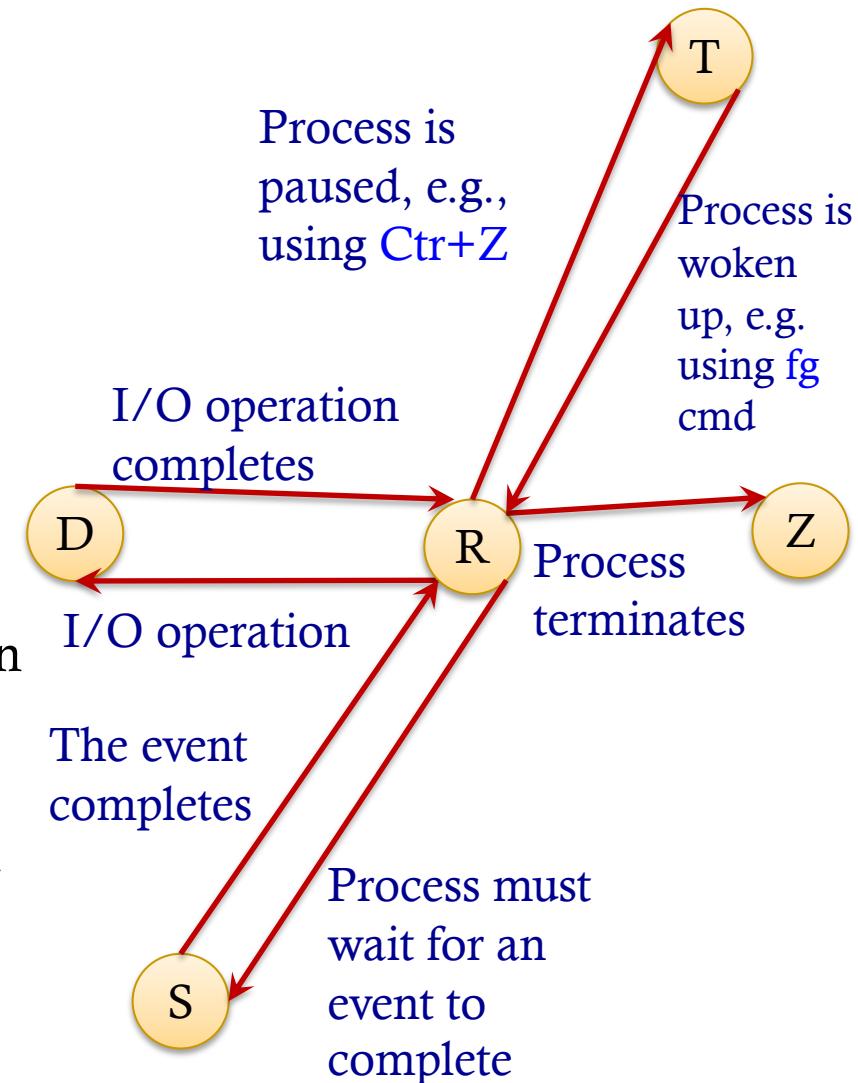
❑ The **states** of a process:

- **New:** The process is being created
- **Ready:** The process is waiting to be assigned to a processor
- **Running:** Instructions are being executed
- **Waiting:** The process is waiting for some event (e.g., I/O completion)
- **Terminated:** The process has finished execution

Process States in Linux

- ❑ Use the `ps` utility to learn process state in Unix:

- Example: `ps -aux`
- Interpreting the output of `ps`
 - **R** the process is running or runnable (on run queue)
 - **D** uninterruptible sleep (usually I/O)
 - **S** interruptible sleep (waiting for an event to complete)
 - **Z** defunct/zombie, terminated but not reaped by its parent (discussed later)
 - **T** stopped, either by a job control signal or because it is being traced



Process Control Block

- ❑ OS uses a Process Control Block (PCB) to represent each process
- ❑ PCB Components:
 - **Process State:** the state of the process, e.g., running, waiting, ...
 - **Process ID:** a unique ID associated with the process, e.g., 1234567
 - **Program Counter:** the address of the next instruction to be executed
 - **CPU Registers:** the current values of the accumulators, stack pointers, etc...
 - **CPU-Scheduling Information:** process priority and other info needed for assigning the process to the CPU

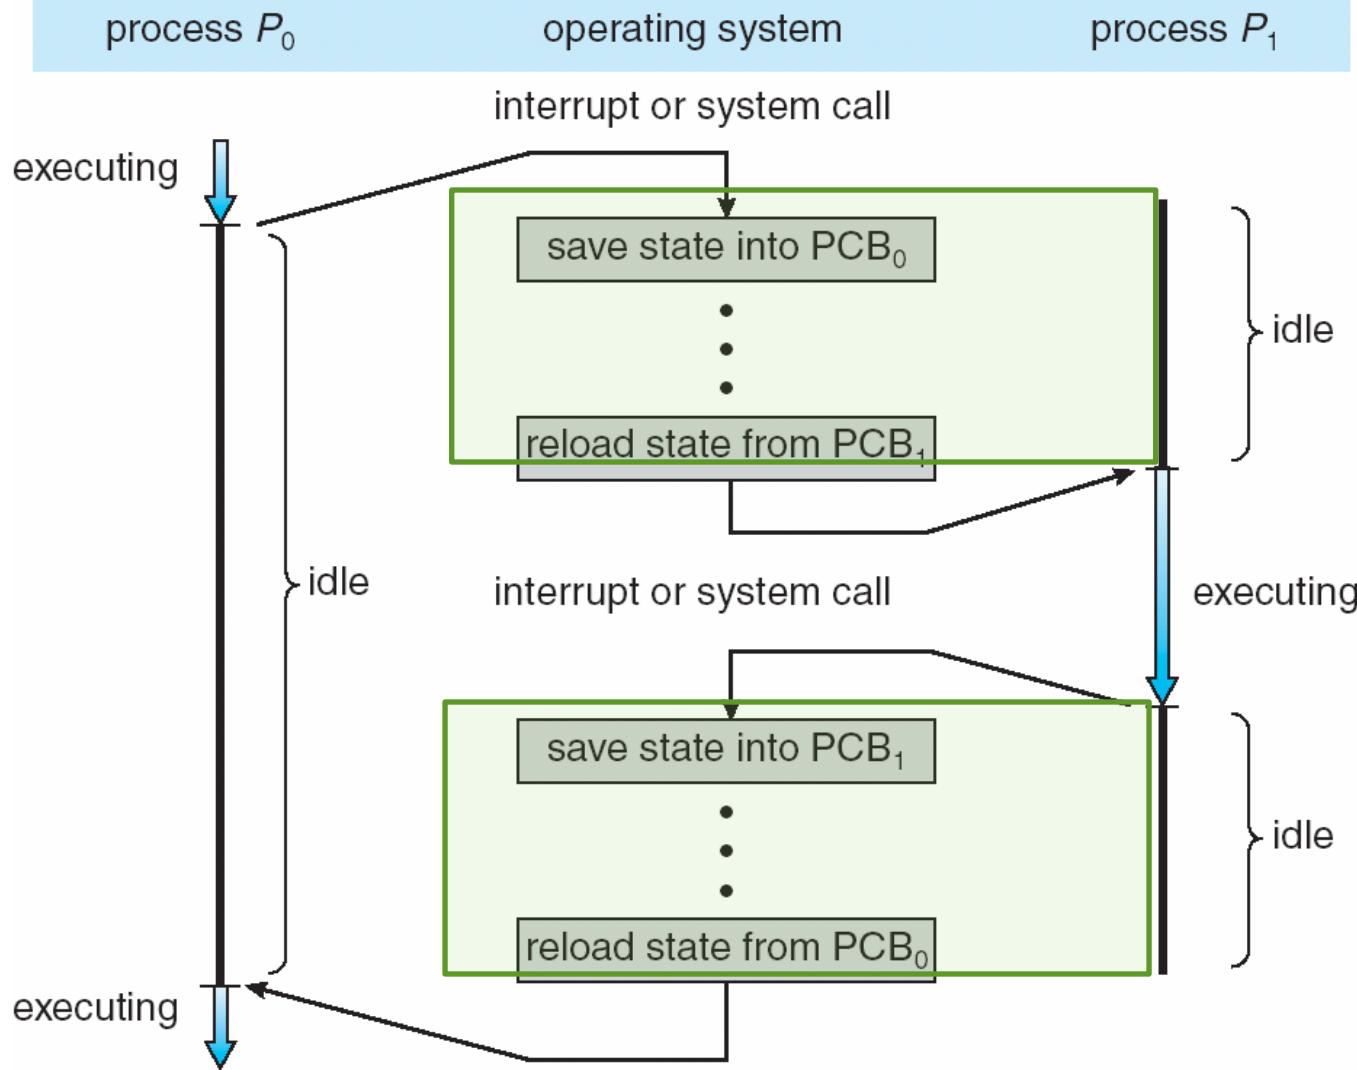


Process Control Block (cont)

- ❑ OS uses a Process Control Block (PCB) to represent each process
 - **Memory Management Information:** info about memory allocated to the process
 - **Accounting Information:** CPU used, clock time elapsed since start (e.g., how long the process was running), time limits, ...
 - **I/O Status Information:** list of I/O devices used by the process, list of open files, etc...



CPU Switch From Process to Process



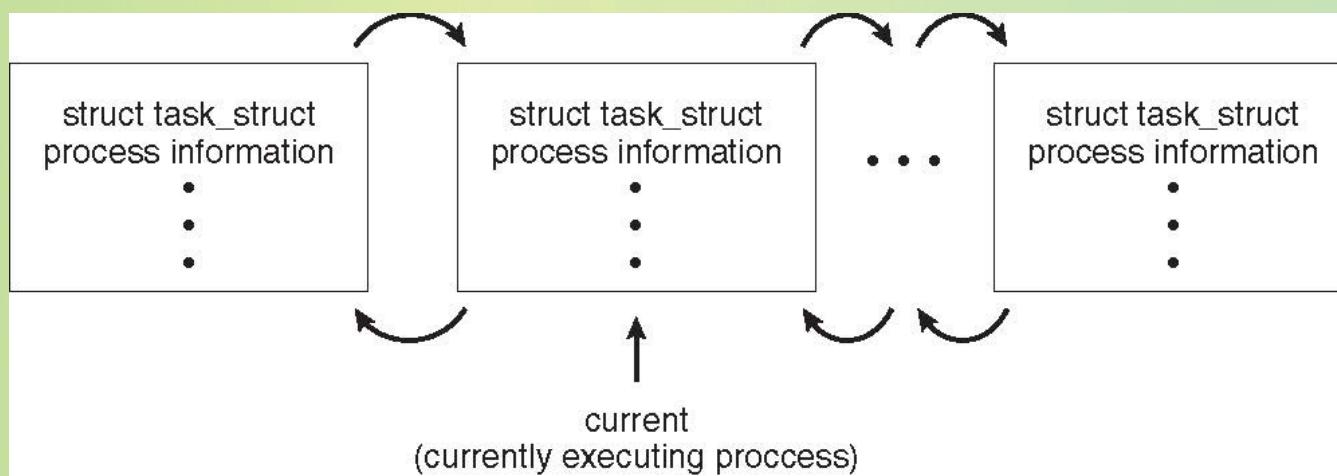
Threads

- So far, process has a single thread of execution
- Consider having multiple program counters per process
 - ◆ Multiple locations can execute at once
 - ▣ Multiple threads of control -> **threads**
- Must then have storage for thread details, multiple program counters in PCB
- See next chapter

Process Representation in Linux

Represented by the C structure `task_struct`

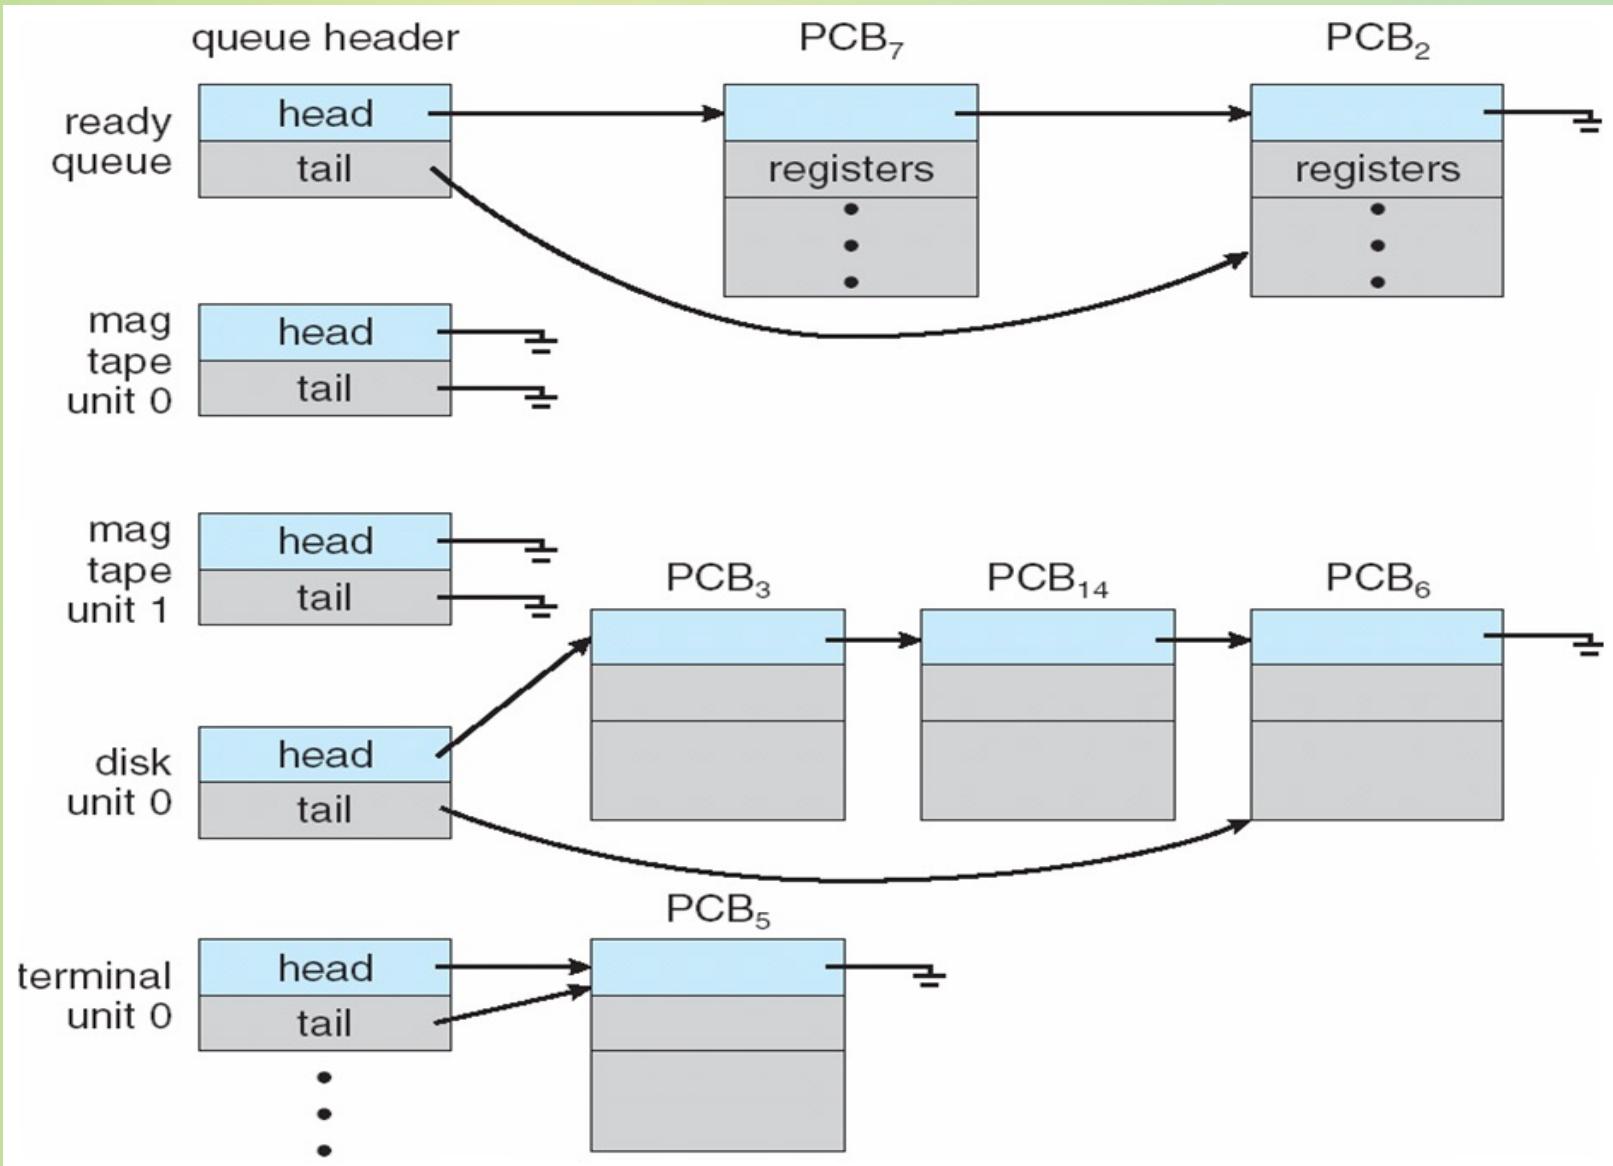
```
pid t_pid; /* process identifier */  
long state; /* state of the process */  
unsigned int time_slice; /* scheduling information */  
struct task_struct *parent; /* this process's parent */  
struct list_head children; /* this process's children */  
struct files_struct *files; /* list of open files */  
struct mm_struct *mm; /* address space of this process */
```



Process Scheduling

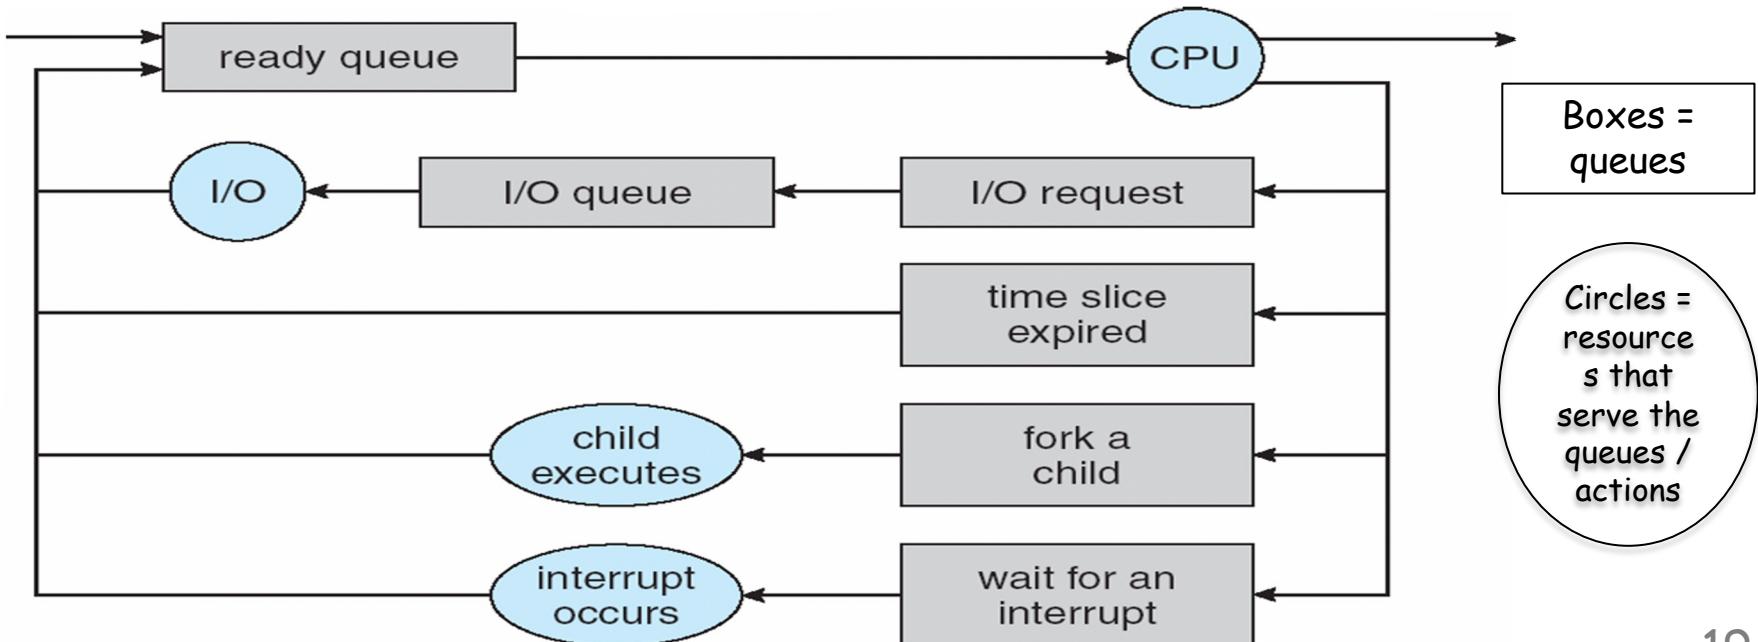
- Maximize CPU use, quickly switch processes onto CPU for time sharing
- **Process scheduler** selects among available processes for next execution on CPU
- Maintains **scheduling queues** of processes
 - **Job queue** - set of all processes in the system
 - **Ready queue** - set of all processes residing in main memory, ready and waiting to execute
 - **Device queues** - set of processes waiting for an I/O device
 - Processes migrate among the various queues

Ready Queue And Various I/O Device Queues



Process Scheduling: The Queuing Mechanism in Action

1. New process is placed on a **ready queue**
2. The process is **dispatched**, i.e., selected for execution
3. During execution:
 - Process may be queued on the device queue due to an **I/O request**
 - Process can create a new **subprocess** and wait for its termination
 - Process can be **interrupted** and removed from the CPU
 - Process **time slice** (i.e., time allotted for its CPU use) expires:
 - Process is removed from the CPU
 - is placed on the ready queue until scheduled to run for another time slice



Schedulers

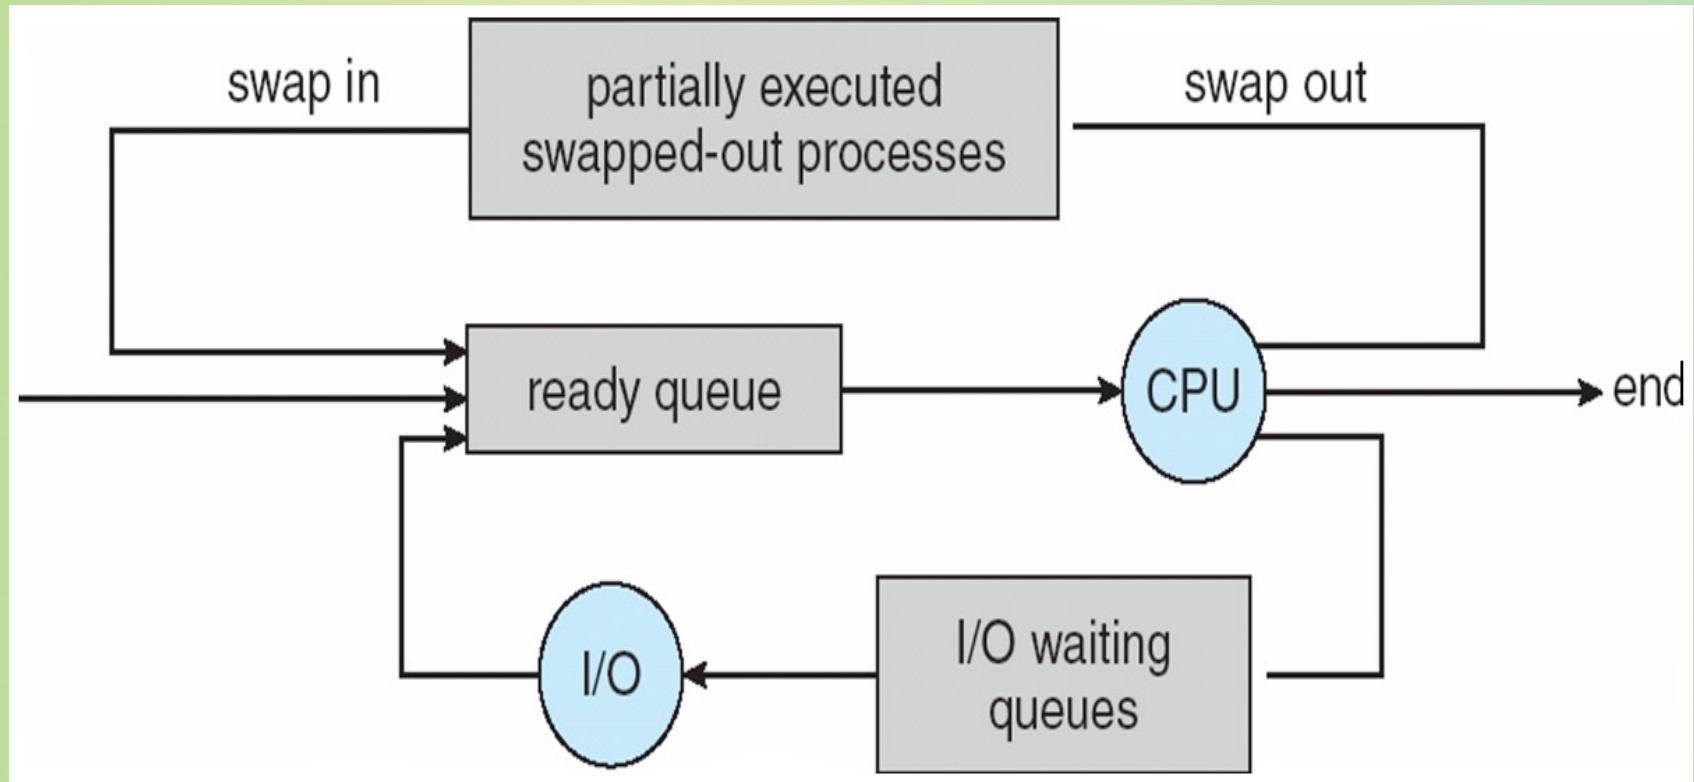
- Short-term scheduler (or CPU scheduler) - selects which process should be executed next and allocates CPU
 - Sometimes the only scheduler in a system
 - Short-term scheduler is invoked frequently (milliseconds) ⇒ (must be fast)
- Long-term scheduler (or job scheduler) - selects which processes should be brought into the ready queue
 - Long-term scheduler is invoked infrequently (seconds, minutes) ⇒ (may be slow)
 - The long-term scheduler controls the degree of multiprogramming
- Processes can be described as either:
 - I/O-bound process - spends more time doing I/O than computations, many short CPU bursts
 - CPU-bound process - spends more time doing computations; few very long CPU bursts
- Long-term scheduler strives for good *process mix*

Process Scheduling: Long-term Schedulers

- What if there are **more** processes than can fit into memory?
 - **Solution:** Temporarily store / spool some processes on a **mass storage** device, e.g., hard drive
- Long-term Scheduler selects spooled processes to load from the mass storage device into main memory
 - Executes **less frequently** than a short-term scheduler
 - **Key idea:** maximize resource utilization by selecting a **mix** of **CPU bound** and **I/O bound** processes
 - **I/O bound processes** are processes that spend more time doing I/O operations than CPU computations
 - **CPU bound processes** are processes that spend more time doing CPU computations than I/O operations
 - **Rational:**
 - If all selected processes are **I/O-bound**, then ready queue will always be empty i.e., nothing to execute on the CPU!
 - If all selected processes are **CPU-bound**, the device queues will be empty i.e., the devices go unused
 - ❖ Not all systems such as UNIX and Windows have long-term scheduler

Addition of Medium Term Scheduling

- **Medium-term scheduler** can be added if degree of multiple programming needs to decrease
 - Remove process from memory, store on disk, bring back in from disk to continue execution: **swapping**



Context Switch

- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process represented in the **PCB**
- Context-switch time is **overhead**; the system does no useful work while switching
 - The more complex the OS and the PCB → the longer the context switch
- Time dependent on hardware support
 - Some hardware provides multiple sets of registers per CPU → multiple contexts loaded at once

Multitasking in Mobile Systems

- Some mobile systems (e.g., early version of iOS) allow only one process to run, others suspended
- Due to screen real estate, user interface limits iOS provides for a
 - Single **foreground** process- controlled via user interface
 - Multiple **background** processes- in memory, running, but not on the display, and with limits
 - Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
- Android runs foreground and background, with fewer limits
 - Background process uses a **service** to perform tasks
 - Service can keep running even if background process is suspended
 - Service has no user interface, small memory use

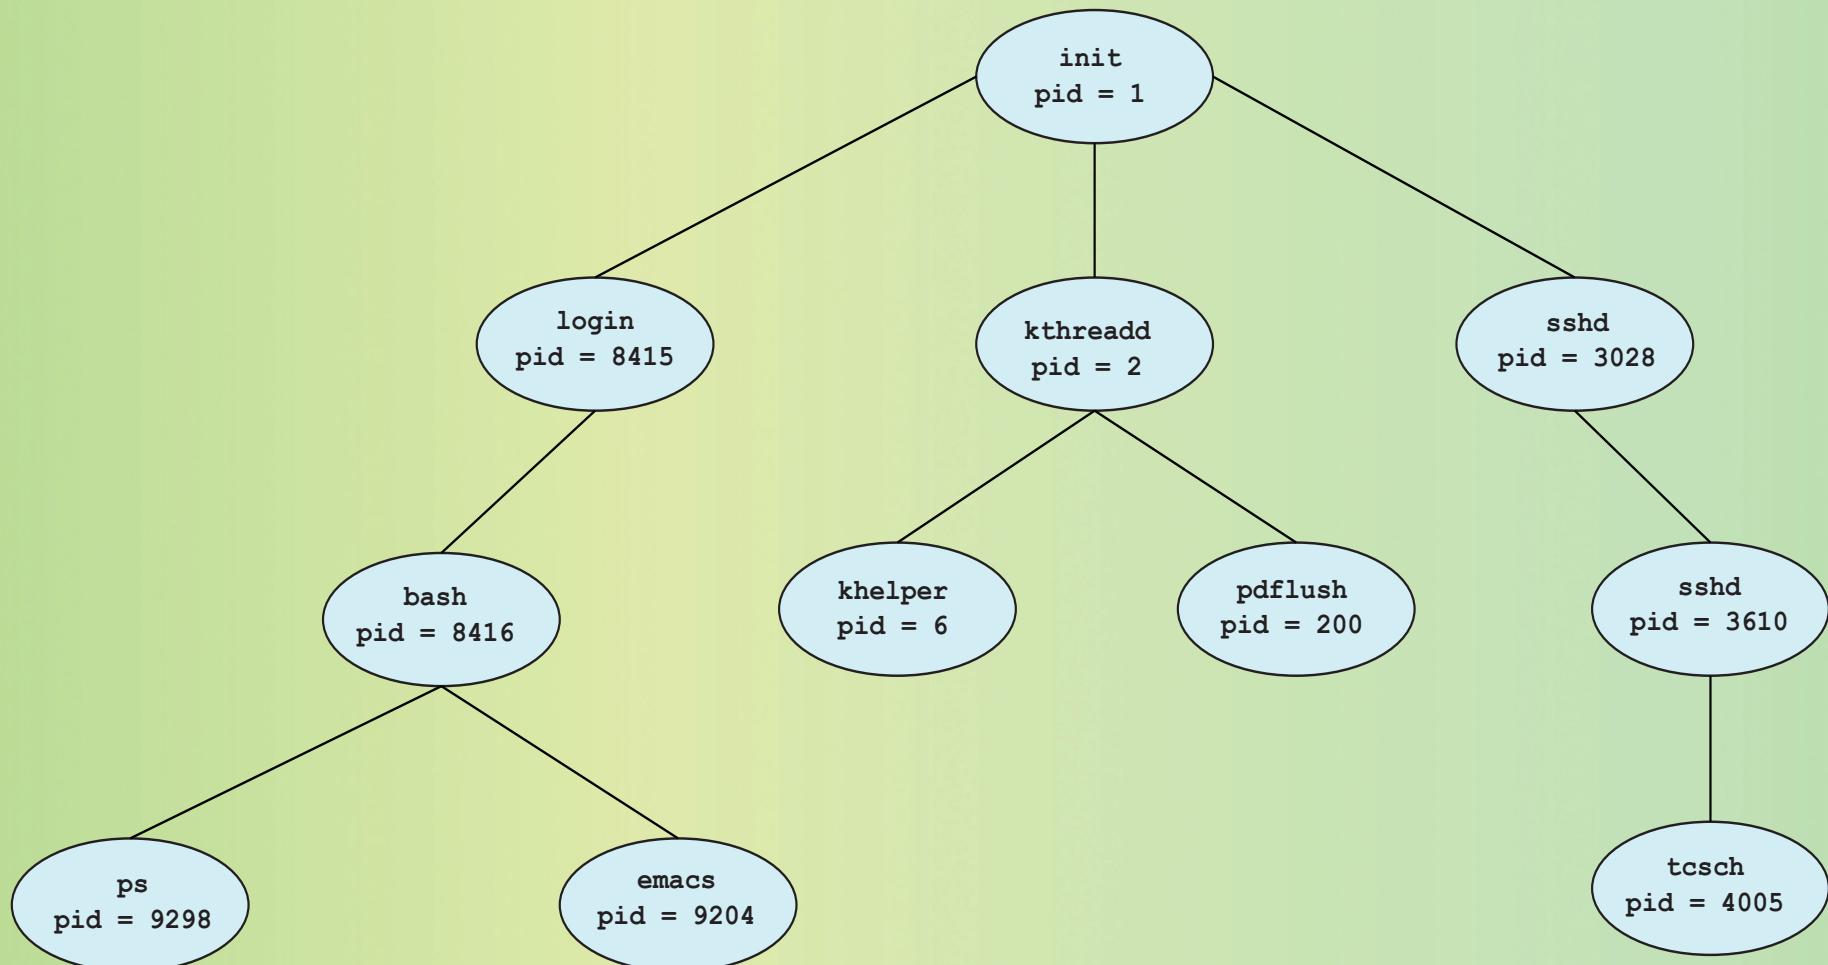
Operations on Processes

- System must provide mechanisms for:
 - ◆ process creation,
 - ◆ process termination,
 - ◆ and so on as detailed next

Process Creation

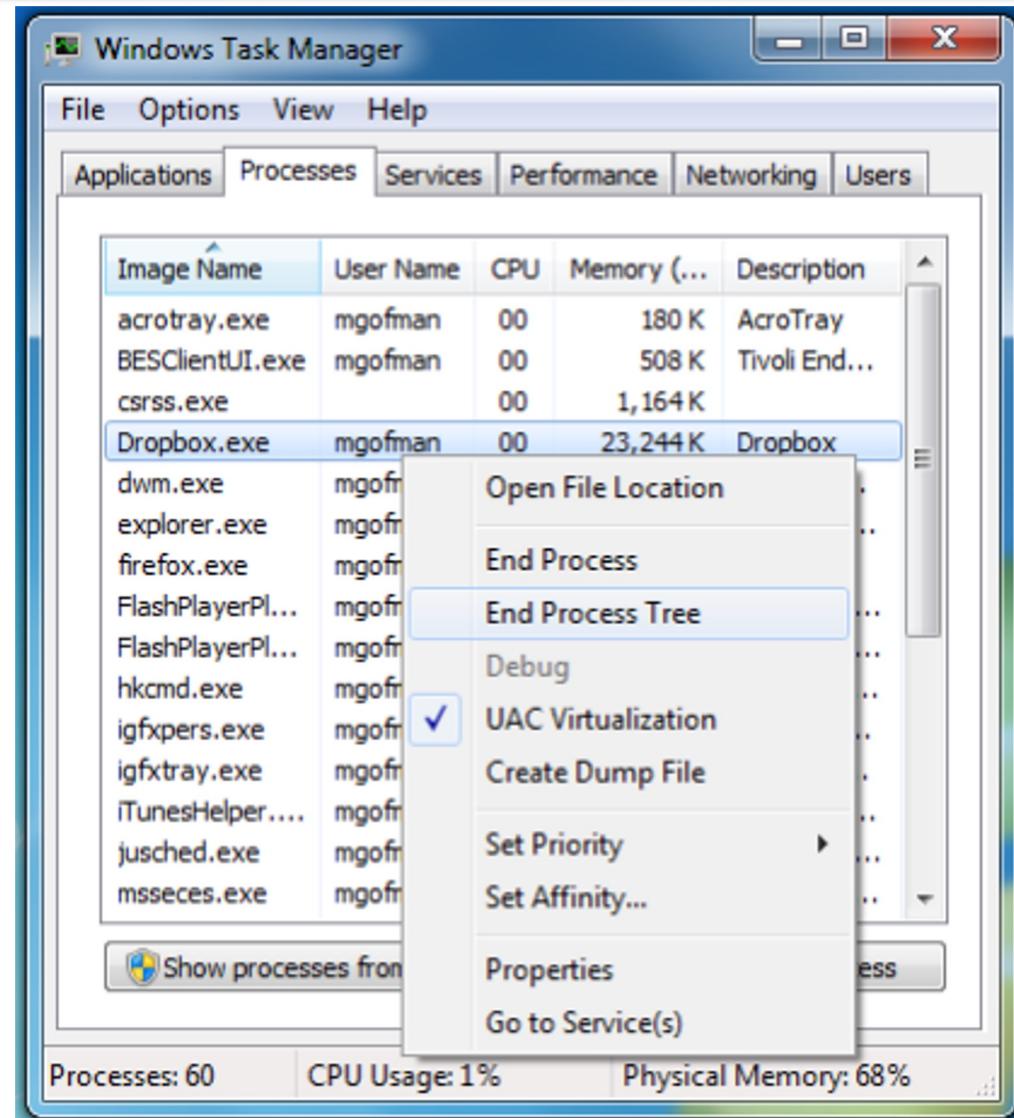
- Parent process create children processes, which, in turn create other processes, forming a tree of processes
- Generally, process identified and managed via a process identifier (pid)
- Resource sharing options
 - ◆ Parent and children share all resources
 - ◆ Children share subset of parent's resources
 - ◆ Parent and child share no resources
- Execution options
 - ◆ Parent and children execute concurrently
 - ◆ Parent waits until children terminate

A Tree of Processes in Linux



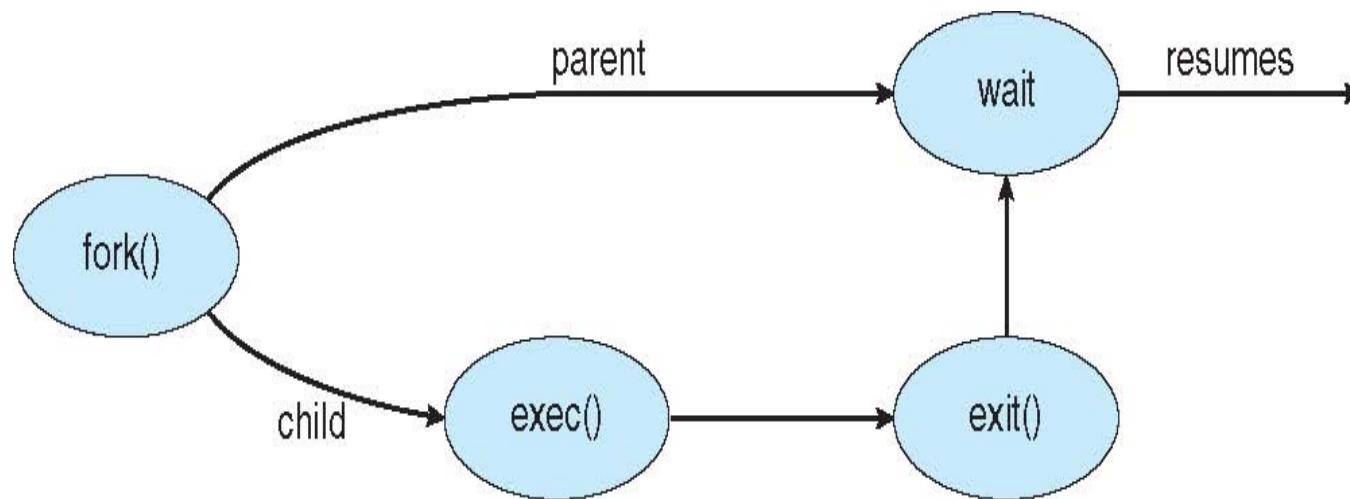
Operations on Processes: Windows Process Tree Example

- Windows: "End Process Tree" terminates the selected process and all its descendants



Unix Process Creation

- ❑ `fork()` system call creates new process
- ❑ `exec()` system call used after a `fork()` to replace the process' memory space with a new program



Process Creation in Unix/Linux: *fork()*

- ❑ **fork()** - a system call is issued by a **parent process** to create a child process
- ❑ Child process is a **clone** of a parent process
- ❑ Both parent and child continue execution at the instruction **immediately after fork()**:
 - In the child **fork()** **returns 0**
 - In the parent **fork** returns **process id (pid)** of the child
 - **fork()** **returns -1** on failure
- ❑ The child process inherits:
 - The set of **files** opened by the parent process
 - Other resources...

Process Creation in Unix/Linux: `exec()/wait()/exit()`

- `exec(...)`- replaces the **program** of the caller process with a **new program**
- `wait(...)`- waits until the child terminates
- `exit(int exitcode)`- terminates the caller process with the specified exit code

Process Creation in Unix/Linux: `exec()` variants

- ❑ `int execl(const char *path, const char *arg, ...);`
- ❑ `int execlp(const char *file, const char *arg, ...);`
- ❑ `int execle(const char *path, const char *arg, ..., char *const envp[]);`
- ❑ `int execv(const char *path, char *const argv[]);`
- ❑ `int execvp(const char *file, char *const argv[]);`
- ❑ `int execvpe(const char *file, char *const argv[], char *const envp[]);`
- ❑ Example: `execlp(const char *file, const char *arg, ...);`
 - `file`: the path of the executable image
 - `arg0...argn`: command line arguments to pass to the process
- ❑ All return `-1` on failure

Process Creation in Unix/Linux: `wait()` variants

- ❑ `wait()` and `waitpid()` return the process id of the child
 - `pid_t wait(int *status);`
 - `pid_t waitpid(pid_t pid, int *status, int options);`

- ❑ `waitid()` return 0 on success and -1 on failure
 - `int waitid(idtype_t idtype, id_t id, siginfo_t *infop, int options);`

Process Creation in Unix/Linux: Putting it all Together

```
// Parent process
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execvp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to
         * complete */
        wait(NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

Start with the parent process

Parent process issues a fork() system call

fork() clones the parent process → Both parent & child continue by executing the next instruction after *fork()*

// Parent process

```
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to
         * complete */
        wait(NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

// Child process

```
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to
         * complete */
        wait(NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

In parent, fork() returns child's process ID

```
// Parent process
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        /*
         * execvp("/bin/ls", "ls", NULL);
        */
    }
    else { /* parent process */
        /* parent will wait for the child to
         * complete */
        wait(NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

In child, fork() returns pid of 0

```
// Child process
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        /*
         * execvp("/bin/ls", "ls", NULL);
        */
    }
    else { /* parent process */
        /* parent will wait for the child to
         * complete */
        wait(NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

Parent issues a `wait(...)` system call to wait until the child terminates

```
// Parent process
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        /*
         * execvp("/bin/ls", "ls", NULL);
        */
    }
    else { /* parent process */
        /* parent will wait for the child to
         * complete */
        wait(NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

Child issues a `execvp()` system call to replace its executable image with that of `ls` command

```
// Child process
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        /*
         * execvp("/bin/ls", "ls", NULL);
        */
    }
    else { /* parent process */
        /* parent will wait for the child to
         * complete */
        wait(NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

Parent waits (in `wait(...)`) for the child process to terminate

```
// Parent process
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to
         * complete */
        wait(NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

`ls` command executes starting from the first instruction; original child code is destroyed

```
// Child process
```

`ls` command code

wait(...) returns, and parent process executes the next

instruction

```
int main()
{
```

```
    pid_t pid;
```

```
    /* fork another process */
```

```
    pid = fork();
```

```
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
```

```
    else if (pid == 0) { /* child process */
        /*
```

```
        execvp("/bin/ls", "ls", NULL);
    }
```

```
    else { /* parent process */
        /* parent will wait for the child to
           complete */
        /*
```

```
        wait(NULL);
    }
```

```
    printf ("Child Complete");
    exit(0);
}
```

ls command finishes execution and terminates

// Child process

ls command code

*Parent process issues an exit()
system call in order to self-*

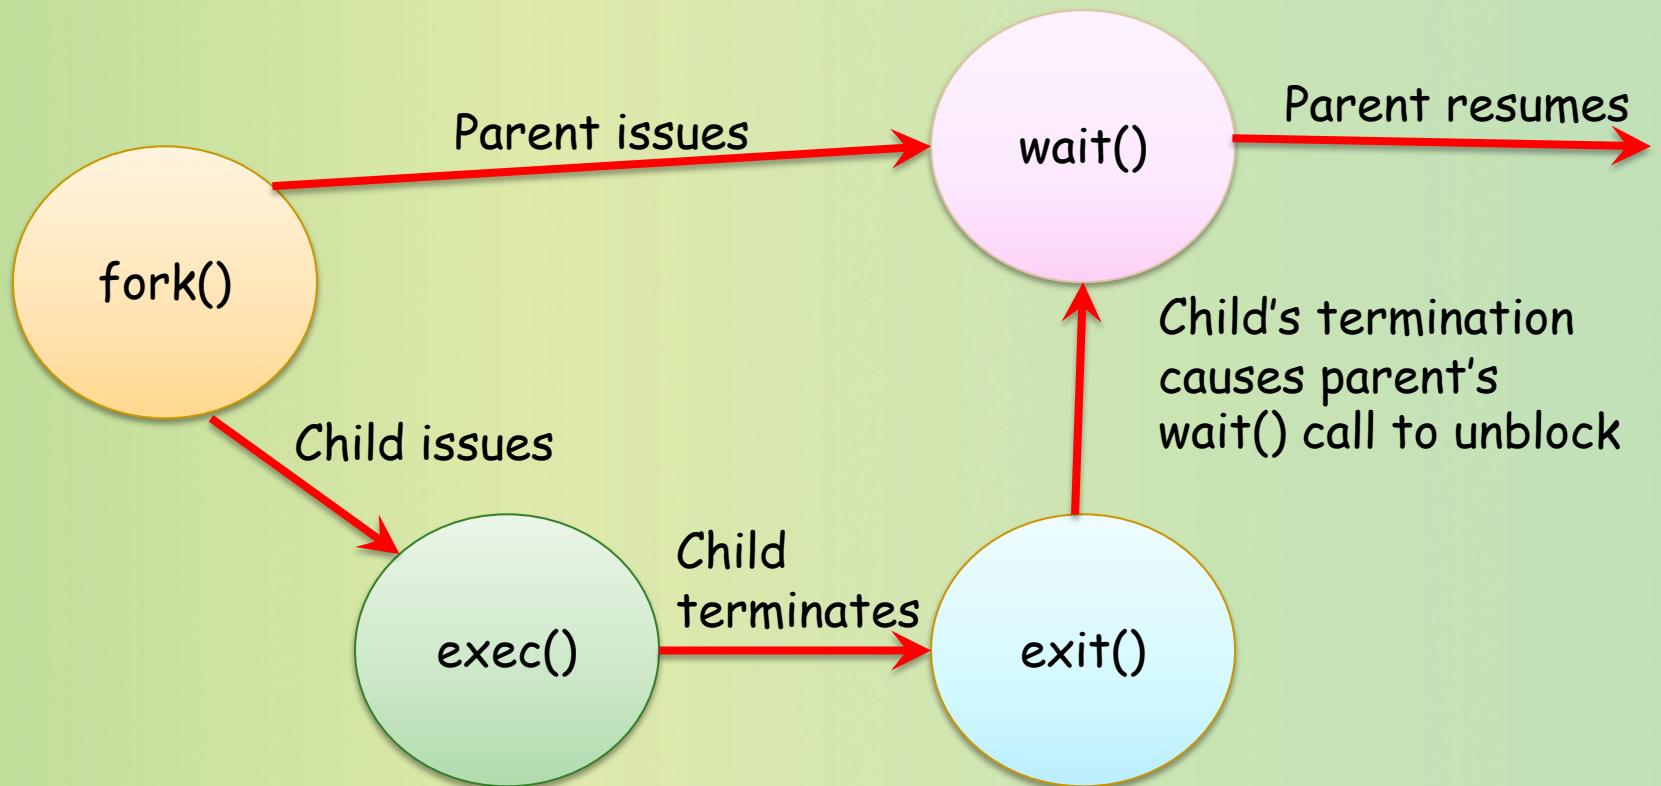
~~terminate~~ process

```
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process
    */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to
         * complete */
        wait(NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

Parent process terminates

```
// Parent process
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execvp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to
         * complete */
        wait(NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

Process Creation in Unix/Linux: Summary of fork()/exec()/wait()



Process Termination

- Process executes last statement and then asks the operating system to **delete** it using the **exit() system call**.
 - ◆ Returns status data from child to parent (via `wait()`)
 - ◆ Process' resources are deallocated by operating system
- Parent may terminate the execution of children processes using the **abort() system call**. Some reasons for doing so:
 - ◆ Child has exceeded allocated resources
 - ◆ Task assigned to child is no longer required
 - ◆ The parent is exiting and the operating systems does not allow a child to continue if its parent terminates

Process Termination

- Some operating systems do not allow child to exists if its parent has terminated. If a process terminates, then all its children must also be terminated.
 - ◆ **cascading termination.** All children, grandchildren, etc. are terminated.
 - ◆ The termination is initiated by the **operating system**.
- The parent process may wait for termination of a child process by using the **wait() system call**. The call returns status information and the pid of the terminated process

```
pid = wait(&status);
```

- If no parent waiting (did not invoke `wait()`) process is a **zombie**
- If parent terminated without invoking `wait`, process is an **orphan**

Zombie Process

- If a parent **forks** a child, but does not issue a **wait()** after the child terminates, the terminated child becomes a **zombie process**
- **Zombie process** - a **terminated** process whose PCB was not **deallocated** - i.e., PCB contains child's **exit code** such as the code returned by int main()
 - The child will remain a zombie until the parent calls **wait()**
 - Child's pid and PCB are released
- Child's **exit code** may be useful to the **parent** - e.g., to see whether the child has exited with an error

Orphan Process

- What if the parent process **terminates** and did not call **wait()** on the child?
 - The child becomes an **orphan process**
 - **init** process becomes the new **parent** of the orphaned children
 - **init** periodically calls **wait()** to collect the return statuses of orphans and release the orphan's process pid and PCB

Multiprocess Architecture - Chrome Browser

- Many web browsers ran as single process (some still do)
 - If one web site causes trouble, entire browser can hang or crash
- Google Chrome Browser is multiprocess with 3 different types of processes:
 - Browser** process manages user interface, disk and network I/O
 - Renderer** process renders web pages, deals with HTML, Javascript. A new renderer created for each website opened
 - Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
 - Plug-in** process for each type of plug-in

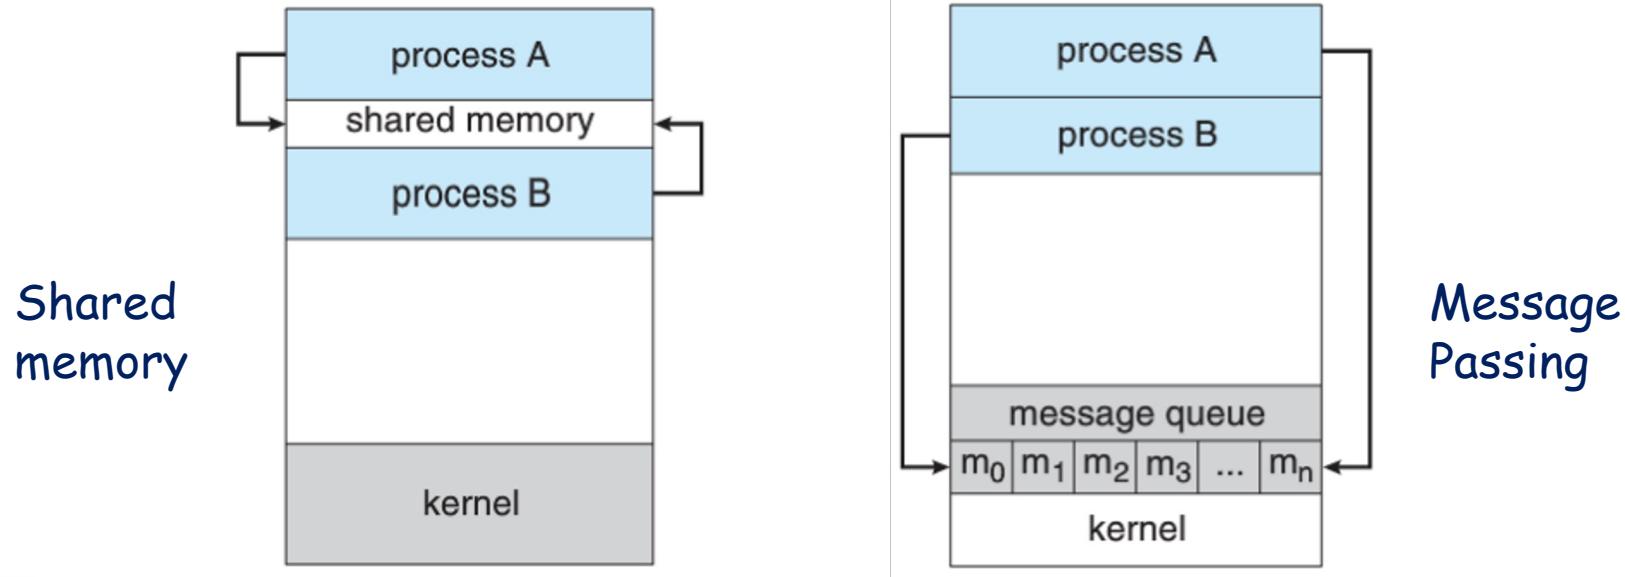


Interprocess Communication

- Processes within a system may be *independent* or *cooperating*
- Cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
 - ◆ Information sharing
 - ◆ Computation speedup
 - ◆ Modularity
 - ◆ Convenience
- Cooperating processes need **interprocess communication (IPC)**
- Two models of IPC
 - ◆ Shared memory
 - ◆ Message passing

IPC Mechanisms for Cooperating Processes

- Shared memory: cooperating processes exchange information by reading/writing data from/to a region of shared memory
- Message passing: cooperating processes exchange messages



IPC: Shared Memory vs Message Passing

❑ Shared memory

- Faster than message passing: only requires intervention from the OS to establish a shared memory region
- Good for large transfers of information
- Disadvantage: requires process synchronization to ensure that no two processes write the same memory location at the same time

❑ Message passing

- No need for synchronization
- Good for small information transfers
- Easier to implement than shared memory
- Disadvantage: usually requires OS intervention on every message transfer
 - Can be slower than shared memory

Producer-Consumer Problem

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
 - ◆ unbounded-buffer places no practical limit on the size of the buffer
 - ◆ bounded-buffer assumes that there is a fixed buffer size

Bounded-Buffer - Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10  
typedef struct {
```

```
    ...
```

```
} item;
```

```
item buffer[BUFFER_SIZE];
```

```
int in = 0;
```

```
int out = 0;
```

Bounded-Buffer - Producer

```
item next_produced;  
while (true) {  
    /* produce an item in next produced */  
    while (((in + 1) % BUFFER_SIZE) == out)  
        ; /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
}
```

Bounded Buffer - Consumer

```
item next_consumed;  
while (true) {  
    while (in == out)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
  
    /* consume the item in next consumed */  
}  
}
```

Interprocess Communication - Shared Memory

- An area of memory shared among the processes that wish to communicate
- The communication is under the control of the users processes not the operating system.
- Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.
- Synchronization is discussed in great details in Chapter 5.

Interprocess Communication - Message Passing

- Mechanism for processes to communicate and to synchronize their actions
- Message system - processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - ◆ *send(message)*
 - ◆ *receive(message)*
- The message size is either fixed or variable

Message Passing (Cont.)

- If processes P and Q wish to communicate, they need to:
 - ◆ Establish a **communication link** between them
 - ◆ Exchange **messages** via send/receive
- Implementation issues:
 - ◆ How are links established?
 - ◆ Can a link be associated with more than two processes?
 - ◆ How many links can there be between every pair of communicating processes?
 - ◆ What is the capacity of a link?
 - ◆ Is the size of a message that the link can accommodate fixed or variable?
 - ◆ Is a link unidirectional or bi-directional?

Message Passing (Cont.)

- Implementation of communication link

- ◆ Physical:

- Shared memory
 - Hardware bus
 - Network

- ◆ Logical:

- Direct or indirect
 - Synchronous or asynchronous
 - Automatic or explicit buffering

Direct Communication

- Processes must name each other explicitly:
 - ◆ **send (P , message)** - send a message to process P
 - ◆ **receive(Q , message)** - receive a message from process Q
- Properties of communication link
 - ◆ Links are established automatically
 - ◆ A link is associated with exactly one pair of communicating processes
 - ◆ Between each pair there exists exactly one link
 - ◆ The link may be unidirectional, but is usually bi-directional

Indirect Communication

- Messages are directed and received from **mailboxes** (also referred to as **ports**)
 - ◆ Each mailbox has a unique id
 - ◆ Processes can communicate only if they share a mailbox
- Properties of communication link
 - ◆ Link established only if processes **share a common mailbox**
 - ◆ A link may be associated with **many processes**
 - ◆ Each pair of processes may share **several communication links**
 - ◆ Link may be **unidirectional or bi-directional**

Indirect Communication

- Operations

- ◆ **create** a new mailbox (port)
- ◆ **send** and **receive** messages through mailbox
- ◆ **destroy** a mailbox

- Primitives are defined as:

send(A, message) - send a message to mailbox A

receive(A, message) - receive a message from
mailbox A

Example: Indirect Communication

- Mailbox sharing
 - ◆ P_1, P_2 , and P_3 share mailbox A
 - ◆ P_1 , sends; P_2 and P_3 receive
 - ◆ Who gets the message?
- Solutions
 - ◆ Allow a link to be associated with at most two processes
 - ◆ Allow only one process at a time to execute a receive operation
 - ◆ Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

Synchronization

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
 - ◆ Blocking send -- the sender is blocked until the message is received
 - ◆ Blocking receive -- the receiver is blocked until a message is available
- **Non-blocking** is considered **asynchronous**
 - ◆ Non-blocking send -- the sender sends the message and continue
 - ◆ Non-blocking receive -- the receiver receives:
 - A valid message, or
 - Null message
- Different combinations possible
 - ◆ If both send and receive are blocking, we have a **rendezvous**

Synchronization (Cont.)

- Producer-consumer becomes trivial

```
message next_produced;  
while (true) {  
    /* produce an item in next produced */  
    send(next_produced);  
}
```

```
message next_consumed;  
while (true) {  
    receive(next_consumed);  
  
    /* consume the item in next consumed */  
}
```

Buffering

- Queue of messages attached to the link.
- implemented in one of three ways
 - 1. Zero capacity - no messages are queued on a link.
Sender must wait for receiver (**rendezvous**)
 - 2. Bounded capacity - finite length of n messages
Sender must wait if link full
 - 3. Unbounded capacity - infinite length
Sender never waits

Address Space

- Process must be able to reference every resource in its abstract machine
- Assign each unit of resource an address
 - ◆ Most addresses are for **memory locations**
 - ◆ Abstract device registers
 - ◆ Mechanisms to manipulate resources
- Addresses used by one process are **inaccessible** to other processes
- Say that each process has its own **address space**

Shared Address Space

- Classic processes sharing program \Rightarrow shared address space support
- Thread model simplifies the problem
 - ◆ All threads in a process **implicitly** use that process's address space , but no "unrelated threads" have access to the address space
 - ◆ Now trivial for threads to share a program and data
 - If you want **sharing**, encode your work **as threads in a process**
 - If you **do not want sharing**, place threads **in separate processes**

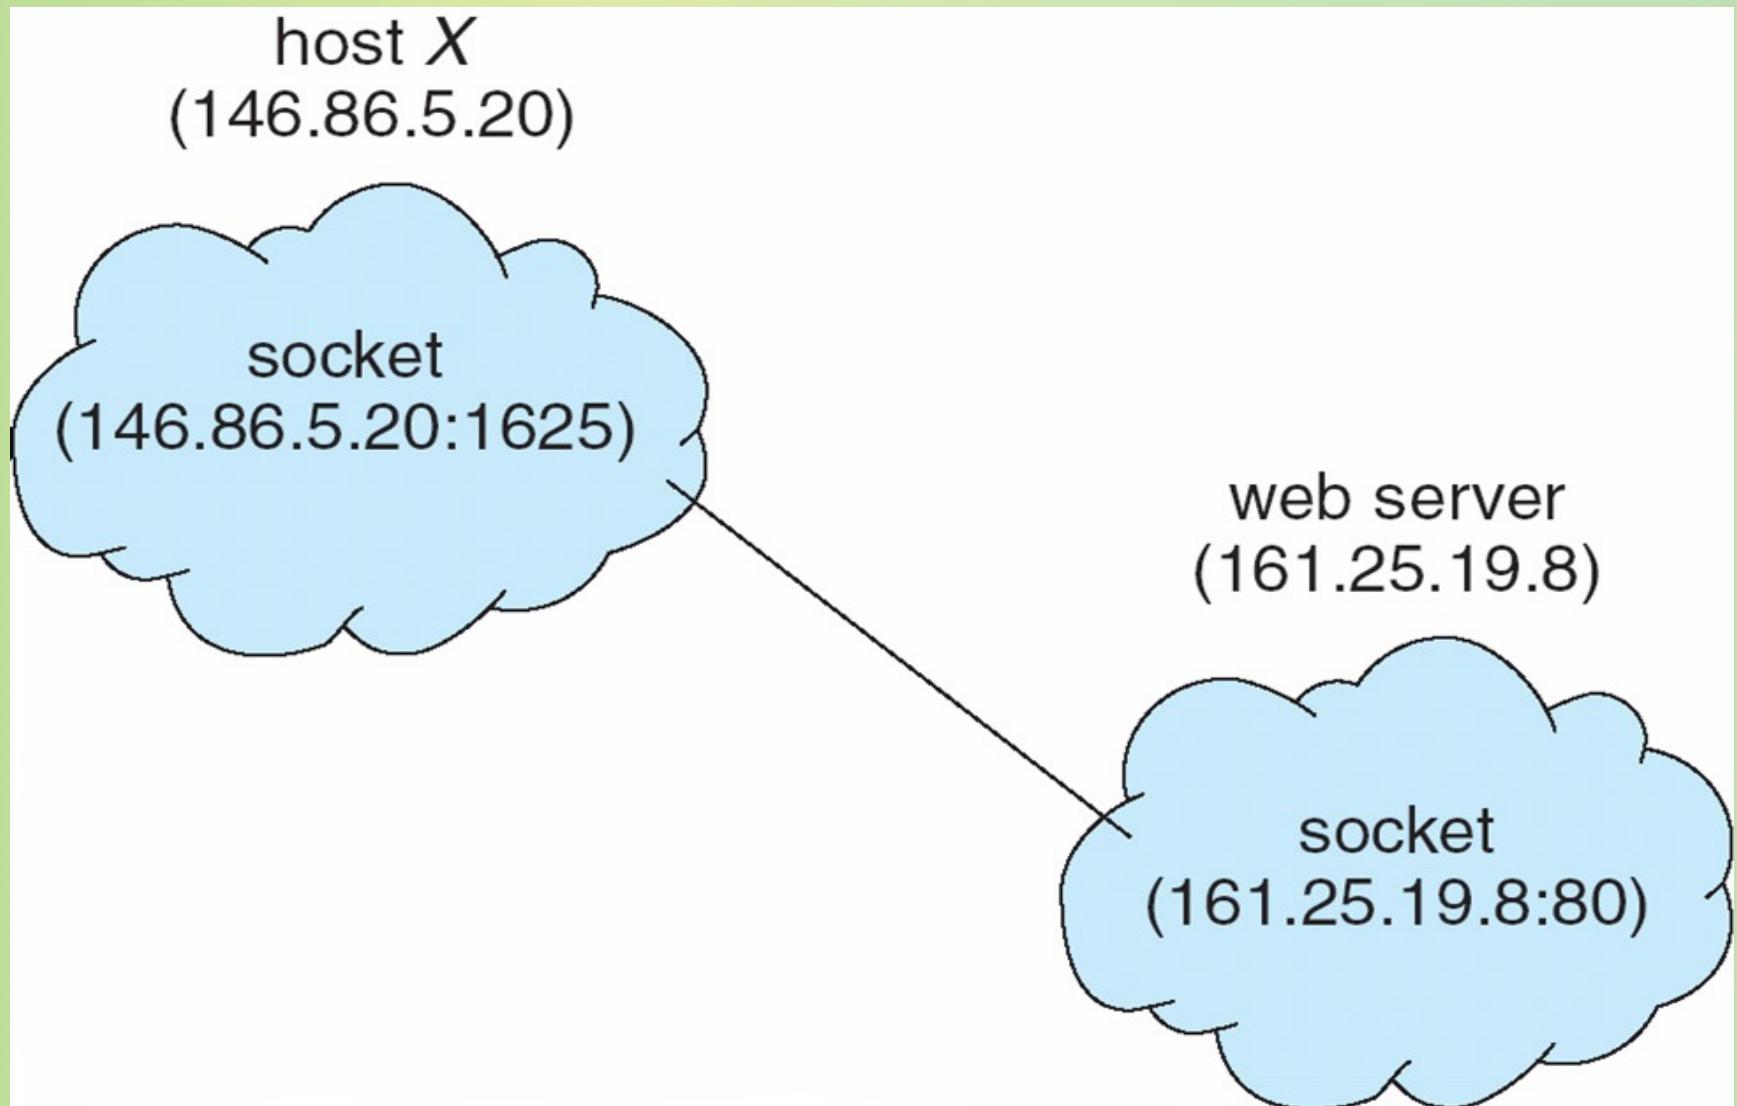
Communications in Client-Server Systems

- Sockets—details are covered in the class of *Computer Communications*
- Remote Procedure Calls
- Pipes

Sockets

- A **socket** is defined as an endpoint for communication
- Concatenation of **IP address** and **port** - a number included at start of message packet to differentiate network services on a host
- The socket **161.25.19.8:1625** refers to port 1625 on host **161.25.19.8**
- Communication consists between a pair of sockets
- All ports below 1024 are **well known**, used for **standard services**
- Special IP address **127.0.0.1** (**loopback**) to refer to system on which process is running

Socket Communication



Remote Procedure Calls

- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems
 - ◆ Again uses ports for service differentiation
- **Stubs** - client-side proxy for the actual procedure **on the server**
- The client-side stub locates the server and **marshalls** the parameters
- The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server
- On Windows, stub code compile from specification written in **Microsoft Interface Definition Language (MIDL)**

Remote Procedure Calls (Cont.)

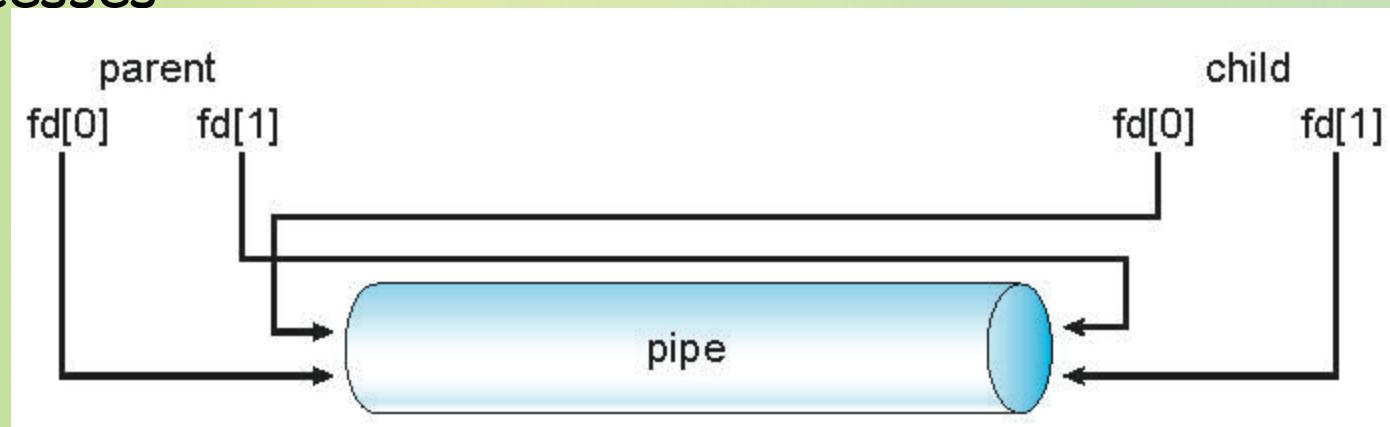
- Data representation handled via **External Data Representation (XDL)** format to account for different architectures
 - ◆ Big-endian and little-endian
- Remote communication has more failure scenarios than local
 - ◆ Messages can be delivered *exactly once* rather than *at most once*
- OS typically provides a rendezvous (or **matchmaker**) service to connect client and server

Pipes

- Acts as a conduit allowing two processes to communicate
- Issues:
 - ◆ Is communication unidirectional or bidirectional?
 - ◆ In the case of two-way communication, is it half or full-duplex?
 - ◆ Must there exist a relationship (i.e., *parent-child*) between the communicating processes?
 - ◆ Can the pipes be used over a network?
- **Ordinary pipes** - cannot be accessed from outside the process that created it. Typically, a parent process creates a pipe and uses it to communicate with a child process.
- **Named pipes** - can be accessed without a parent-child relationship.

Ordinary Pipes

- Ordinary Pipes allow communication in standard producer-consumer style
- Producer writes to one end (the write-end of the pipe)
- Consumer reads from the other end (the read-end of the pipe)
- Ordinary pipes are therefore unidirectional
- Require parent-child relationship between communicating processes



- Windows calls these anonymous pipes
- See Unix and Windows code samples in textbook

Named Pipes

- Named Pipes are more powerful than ordinary pipes
- Communication is **bidirectional**
- **No** parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems